PANTA V20 UX Research and Validation Plan

PANTA Product and Engineering

27 August 2026

Contents

1 Research objective **1**

2 Primary hypotheses **1**

3 Participants **1**

4 Core tasks **2**

4.1 Task 1 - Transcript claim . . . 2
4.2 Task 2 - Coverage discrepancy . . . 2
4.3 Task 3 - Derivation . . . 2
4.4 Task 4 - Mission . . . 2
4.5 Task 5 - Spine change . . . 2
4.6 Task 6 - Financing scenario . . . 2
4.7 Task 7 - Decision and settlement . . . 2

5 Quantitative measures **2**

6 Qualitative probes **3**

7 Validation gates for V21 **3**

8 Known research limits **3**

1 Research objective

Determine whether PANTA V20 helps venture professionals form, challenge and preserve a sparse-evidence investment case without mistaking polished AI output for institutional truth.

2 Primary hypotheses

1. Professionals can distinguish an observed speech act from the truth of its content.
2. Candidate discrepancies surface non-obvious conflicts without being mistaken for resolved conclusions.
3. Deterministic derivations are trusted because inputs, formula, units and assumptions are inspectable.
4. AI hypotheses are understood as proposals and are not propagated prematurely.
5. The governed question spine fits venture work without becoming a rigid checklist.
6. Missions improve next-best-work selection while preserving data-egress and authority boundaries.
7. Venture financing and technical validation fit the same product architecture as buyout and growth.
8. Lens changes are understood as institutional policy projections rather than fact changes.

3 Participants

Recruit 12-18 participants across:

- venture associates and principals;
- deep-tech or defence investors;
- technical diligence specialists;
- investment partners/IC members;
- legal/compliance or information-security reviewers;
- portfolio/follow-on investors.

Include at least four users who routinely work from calls and customer references rather than mature data rooms.

4 Core tasks

4.1 Task 1 - Transcript claim

Open a customer utterance and explain what PANTA knows versus what remains asserted.

Success: participant distinguishes statement occurrence from domain truth without prompting.

4.2 Task 2 - Coverage discrepancy

Review the 800 m claim, controlled result and pilot-implied radius.

Success: participant identifies dimensions that must be aligned before calling the values contradictory.

4.3 Task 3 - Derivation

Validate the implied-radius calculation.

Success: participant finds inputs, formula, assumptions, units and method hash in two gestures.

4.4 Task 4 - Mission

Choose how to close the most critical technical unknown.

Success: participant can distinguish internal analysis, public research, human outreach and physical test authority classes.

4.5 Task 5 - Spine change

Assess whether maintenance economics deserves a standalone question.

Success: participant understands affected bindings, historical alias preservation and authority requirement.

4.6 Task 6 - Financing scenario

Compare base and six-month-delay paths.

Success: participant identifies runway, dilution, next-round need and follow-on reserve without assuming Working changed Current.

4.7 Task 7 - Decision and settlement

Attempt to settle the coverage-restatement scenario without authority.

Success: the user predicts the block, understands the Human Stop and completes the governed path.

5 Quantitative measures

- time to locate exact utterance;
- percentage correctly identifying content class;
- time to explain discrepancy dimensions;
- derivation verification success;

- mission-classification accuracy;
- incorrect assumption that prepared equals executed;
- incorrect assumption that Lens changes facts;
- time to understand blocked partial settlement;
- number of gestures from question to exact source;
- perceived confidence in replaying a past decision.

6 Qualitative probes

- Which part would you trust enough to cite in IC?
- Which part still feels like AI theatre?
- What would cause you to reject a generated discrepancy?
- When should a new question require Partner approval?
- What confidential context would you refuse to send to an external research service?
- Which technical performance dimensions are missing?
- Does the cap-table/runway model match your real venture workflow?

7 Validation gates for V21

V21 should not claim production venture readiness until:

- at least 80% of participants correctly distinguish speech occurrence from content truth;
- at least 80% verify one derivation without facilitator help;
- no participant mistakes a synthetic mission for live research after reading the disclosure;
- no authority holder can settle the Human Stop through the UI without the server record;
- at least two genuinely different venture sectors render without component forks;
- practitioners validate the venture question archetype and identify accepted optionality.

8 Known research limits

The Tethys fixture is synthetic. It tests product structure and interaction, not the semantic accuracy of any real deal. The packaged runtime does not test live meeting ingestion, real SSO/RBAC, production egress controls, real external research, human outreach or physical-test orchestration.